

AI for Enterprise

Powering Intelligence

MARCH 2026



“Maybe there's a job in the world that AI won't change, but I haven't thought of it.”

– Doug McMillon, CEO, Walmart

“AI could wipe out half of all entry-level white-collar jobs”

– Dario Amodei, Anthropic CEO

"AI will make a few people much richer and most people poorer."

– Geoffrey Hinton

"The real impact of AI doing jobs in the next few years will begin to be palpable."

– Sam Altman, CEO, OpenAI

“[AI]... will make us smarter and more efficient for a time, and will unlock enormous amounts of economic growth, but [is] fundamentally labor replacing.

– Mustafa Suleyman, CEO, Microsoft AI

“I can easily imagine a world where 30 to 40% of the tasks that happen in the economy today get done by AI in the not very distant future ”

– Sam Altman, CEO, OpenAI

“The clock is ticking, and they’re keeping score.”

– Jester, Top Gun

“Artificial intelligence is going to replace literally half of all white-collar workers."

– Jim Farley, CEO, Ford

"In less than 20 years, working at all will be optional."

– Elon Musk, CEO, Tesla & SpaceX

"You're not going to lose your job to AI. You're going to lose your job to someone who uses AI. Every job will be affected, and immediately."

– Aravind Srinivas, CEO, Perplexity AI

"We're going to see AI having the capabilities to replace many, many jobs."

– Geoffrey Hinton, Nobel Laureate, "Godfather of AI"

AI will change the future of work.

Empowering employees,
transforming operations,
and shortening the path from
data to insight to action.

30%

of US work hours
automated by AI by 2030¹

1.1B

Jobs transformed by AI
over the next decade²

Source: ¹ McKinsey, *Generative AI and the future of work in America*, Jul 2023

² World Economic Forum, *The Reskilling Revolution*, Jan 2026

AI adoption is nearly universal. Value realization is not.

ADOPTION

88%

of enterprises use AI in at least one function

2x

Production deployments doubled YoY

Source: McKinsey (Nov 2025), ISG (2025)

INVESTMENT

\$37B

enterprise AI spending tripled in 2025

89%

of organizations expect higher AI spending in 2026

Source: Menlo Ventures (Jan2026), ETR (Nov 2025)

THE GAP

34%

of enterprises are using AI to deeply transform their business

6%


of organizations are "AI high performers" seeing 5%+ EBIT impact

Source: Deloitte (2026), McKinsey (Nov2025)

You know AI is the
future. How do
you create real
business value
today?


We see two dominant enterprise patterns

The answer isn't to wait or to sprint. It's to **commit, measure, and persist.**



Paralysis

Waiting for the landscape to settle before committing.



Scatter

Chasing dozens of tactical wins without a coherent strategy for how they add up.

When it works, it really works

Organizations seeing results from AI aren't slowing down — they're accelerating investment, scaling personalization, and putting real budget behind it.

Across Merkle's research of 100 US c-suite leaders, the organizations getting AI right share common patterns. The differences in outcomes are significant.

84%

grew revenue last year

They're reinvesting, not retreating.

- Only 6% saw any decline
- 89% expect tech budgets to increase further

81%

say GenAI is a priority or significant investment

The money is following the conviction

- Near-universal commitment, only 1% say they won't invest at all
- Top priority: use AI to cut costs without cutting quality

87%

have deployed or piloting AI-powered personalization

Clients are starting to mature past the pilot phase

- 42% say AI is already central to their CX personalization
- Another 45% are actively testing it, not just talking about it


Success patterns are emerging

What do winners do differently?

Five patterns from research and our delivery experience

Across Merkle's research, McKinsey's analysis of **200+ transformations**, and our delivery work, the organizations that escape pilot purgatory share five patterns.

These aren't about technology choice. They're about how organizations approach AI.





01 | FOLLOW THE DATA GRAVITY, NOT THE HYPE

Follow the gravity, not the excitement

When organizations ask "**where should we start with AI?**", the answer is almost always closer than they think.

Your data already clusters around specific systems and processes. Those clusters tell you where AI will generate value fastest.

Three questions a leader can ask to identify their data gravity



Where do your teams already go to answer business questions today?



Which systems contain the data your highest-value decisions depend on?




Where is your data most complete, governed, and trusted?

The answer to all three questions is often **the same system. Start there.**



The enterprise AI data gap

Most enterprise data was designed for humans. AI needs something **fundamentally different**.



Two shifts that close the gap

01

Curate the data

Purpose-built datasets scoped to specific use cases, combining structured, unstructured, and document data into a unified, queryable layer. Documented business logic.

02


Describe the data

A semantic layer that encodes what the data means, how it relates, and how the business uses it. The translation layer between raw data and AI reasoning.



02 | START WITH THE DECISION, NOT THE TECHNOLOGY

Specific use cases beat
enterprise-wide
transformation programs





Specific use cases beat enterprise-wide transformation programs

A well-defined use case answers three questions:

01

What decision does this enable?

02

What delay does it eliminate?

03

How will you know in 6 months?

If you can't answer all three, the use case isn't ready.

This means saying no to the 47 other AI ideas floating around. It means defining success in terms a CFO would recognize, not in terms of "AI maturity."



Defining success before you build

The rigor matters more than the precision of the numbers



Efficiency

- Time to answer
- Throughput

Measured in weeks



Quality

- Accuracy
- Consistency
- Trust

Measured in months



Capability

- Questions you couldn't answer before

Measurable immediately



Adoption

- Who's using it
- Satisfaction

Early signals in weeks

Pick one or two metrics per dimension.

You're building a hypothesis you can defend when leadership asks: *"How will we know this worked?"*



03 | REDESIGN THE WORK, NOT JUST THE TOOLS

The operating model matters more than the model

You can't AI your way out of a poor workflow

AI high performers are **nearly 3x as likely** to have fundamentally redesigned workflows. Workflow redesign is one of the **strongest predictors** of enterprise-level EBIT impact.

Yet education remains the **#1** talent response. Redesign barely registers.

64%

say competing transformation initiatives are the top change barrier

76%

report fragmented processes as a significant work design challenge

AI doesn't fix broken processes. It amplifies them.

Sources: McKinsey (Nov 2025), Merkle "Re-wiring the enterprise", Dec 2025



03 | REDESIGN THE WORK, NOT JUST THE TOOLS

Five recommendations from 100 C-suite conversations

- 01 Establish cross-functional alignment.**
Leading orgs are 2.5x more likely to maintain a unified vision across marketing, tech, and data
- 02 Formalize change management.**
40% of developing organizations rely on informal approaches. This can't be a side project
- 03 Close the skills gap with training AND workflow redesign.**
Education addresses skills, workflow redesign addresses structure, you need both
- 04 Invest in work design.**
Redefine roles, codify human-AI collaboration, implement cross-functional operating models
- 05 Align metrics to business value**
Track revenue impact and forecasting accuracy, not vanity metrics.



Trust must be engineered, not assumed

AI systems are probabilistic. Same question, different answers. Some will be wrong. Fortunately, with the right evaluation framework and iterative improvements, this can be directly addressed.

✕ Without Measurement

Teams check everything manually, eliminating efficiency gains

Executives can't quantify ROI or justify investment

The organization drifts back to the old way

✓ With Measurement

Accuracy is baselined, failures are surfaced, improvement is continuous

Performance is tracked against defined quality standards

Trust compounds over time

83% of AI leaders report major or extreme concern about generative AI progress, an **8x increase** in two years.

- Lucidworks 2025



Getting to production isn't the hard part. Sustaining value is.

Your AI systems are a reflection of your organization. Both are living and breathing.

Monitor from day one

- Accuracy
- Usage
- Failure patterns

Build feedback loops

Users inform system refinement

Budget for improvement

Continuous improvement is TCO, not after thought

Build where you are

Client-owned infrastructure they can evolve

Accept “done” doesn’t exist

A system at 85% today is 70% in six months if nobody’s watching

Two modes of intelligence. You need both.

Most organizations operate entirely in pull mode, where people go find the answers they need. The shift isn't to eliminate that. It's to add a proactive layer that delivers routine intelligence automatically, so your people spend their time on the decisions that require human judgment.



Pull (Reactive) : The **user** initiates

MECHANISMS

Dashboards | Conversational interfaces |
Ad-hoc queries

STRENGTH

Flexible, exploratory, good for novel questions and
deep investigation

LIMITATION

Requires someone to know what to ask and when to
ask it

BEST FOR

- Exploratory analysis
- Root cause investigation
- Complex or ambiguous questions

Human Judgment Layer

Analyst review on critical decisions
before action

Context and interpretation AI can't
provide

Investigation of anomalies flagged
by proactive systems

Sign-off on high-stakes
recommendations

MECHANISMS

Scheduled briefings | Triggered alerts |
Automated digests

STRENGTH

Consistent, timely, doesn't depend on someone
remembering to check

LIMITATION

Only covers what you've anticipated and configured

BEST FOR

- Routine monitoring
- Threshold-based alerts
- Recurring reporting


Push (Proactive): The **system** initiates



AI handles the intelligence. Your people handle the judgment.

Better decisions, at every altitude, every day

The organizations creating the most value from AI aren't building one dashboard or one chatbot. They're building intelligence systems that serve every level of the business, delivering what each altitude needs, in the format and frequency the decision requires.



The goal isn't more data at every level. It's the right intelligence, at the right altitude, at the right moment.

From patterns to practice

How Merkle Helps

Most firms speak the CMO's language or the CIO's. Merkle is **natively bilingual**.



Your AI is only as smart as your data layer

Data Foundations


- AI readiness assessment
- Data architecture and governance
- Structured + unstructured integration



Start with the decision, not tech.
Redesign the work, not just the tools

Work Design

- Use case prioritization
- Operating model transformation
- Workflow redesign and change management



Intelligence at every altitude

Enterprise Intelligence Systems

- Conversational insights systems
- Proactive intelligence delivery
- Automated reporting



Earn trust through measurement

AI Assurance & Trust

- Evaluation framework design & implementation
- Continuous improvement & degradation tracking
- Value measurement tied to business outcomes

Start Here

We start every AI for Enterprise engagement with a structured workshop designed to cut through the noise.

THOUGHT LEADERSHIP SESSION

45-60 MINUTES

- Framework
- Research
- Spark conversation

HALF-DAY WORKSHOP

2-4 HOURS

- AI readiness assessment
- Use case prioritization
- Initial roadmap

MULTI-DAY ENGAGEMENT

2-3 DAYS

- Stakeholder interviews
- Data assessment
- Workflow analysis
- Prioritized roadmap

You will walk away with



AI readiness assessment



Prioritized opportunity
roadmap



Workforce and operating
model implications



Executive alignment on
priorities



Governance and evaluation
recommendations

Case Studies

Case studies



Home Improvement **AI Automation**

Merkle partnered with major Paint Retailer to integrate Palantir into their Microsoft D365 CRM, dramatically streamlining previous speed-to-market bottlenecks. Sherwin-Williams' traditional MDM solution left critical datasets undelivered for months. Palantir closed that gap in weeks and transformed their testing workflow, cutting testing time by 80% without compromising quality or accuracy.



QSR

AI-Empowered Employee Application

A major restaurant brand's AI chatbot was returning inaccurate results across complex business queries. We built a retrieval optimization layer and evaluation framework that lifted agent accuracy from 58% to 90%+, unlocking confidence to scale to 100k+ users.

Building a personalization capability across a complex CX organization

BUSINESS CHALLENGE

Our client struggled to align its data-driven Experience, Marcom, eCommerce, Segment Marketing, In-store Experience, and Research teams, preventing the organization from delivering the omni-channel, personalized customer experiences needed to drive growth.

MERKLE SOLUTION

Gaps identified causing prevention of an agile planning process across disciplines. Built framework for integrated goals and process. Provided methodology and tools for delivering customer driven marketing plans.

- 19 personalization opportunities identified for the residential repaint market segment
- 10 personalization opportunities prioritized based on LOI / LOE scoring
- 3 executive 1:1 CX transformation sessions conducted with the S-W VP of customer insights and experience






Cochlear®

Elevating product-first marketing to **customer-first experiences**



HALEON

Accelerating transformation with **new operating model** constructs



T.Rowe Price®

Creating a client-centric organization through **agile transformation**



Walgreens

Unlocking **marketing value** through technology and operational agility

Growth support

What to look out for

If you hear a client say any of these, **that's a conversation for this team.**

When you hear it, **call us.**
We'll **help you qualify and shape the opportunity.**



Data Readiness

- *"How do we ensure our data is AI-ready?"*
- *"We have lots of data that is under-utilized but think it could be valuable for AI."*
- *"Implementation of best practices for semantic layer architecture."*
- *"Which data to prioritize for our AI roadmap?"*
- *"Right mix of structured and unstructured data for our use case."*



Organizational Alignment

- *"What should humans do vs. the system?"*
- *"Every department has a different AI vision."*
- *"Training isn't changing how we actually work."*
- *"Don't know which use cases to prioritize?"*
- *"Struggling to quantify value of AI investments."*



AI Systems & Trust

- *"Where should we build our agents, and on what?"*
- *"We need support using AI to streamline or automate manual workflows."*
- *"Many AI pilots, no idea which to scale..."*
- *"No framework to measure accuracy or quality."*
- *"Promising AI initiatives but struggling to scale from pilot to production."*



Intelligence & Insights

- *"Nobody looks at our dashboards."*
- *"Two weeks to answer a basic question."*
- *"Insights arrive after the decision is made."*
- *"Our team is spending too much time on manual reports and readouts"*
- *"Drowning in data, starving for answers."*
- *"We don't see to be able to identify opportunities and risks before it's too late."*

Marketing supports

AVAILABLE NOW: Downloadable Infographic:
The Transformation Gap: Why Enterprise AI Investment Isn't Paying Off
Available on Merkle.com

MARCH

The \$2T Question: Why do
95% of AI Investments Fail?
Thought Leadership

MAY

The “People Stack”:
Transformation Means More
than Tech
Thought Leadership

JUNE


Building the Intelligent
Analytics Workspace with
Snowflake
Thought Leadership, Partners

AVAILABLE NOW

INFOGRAPHIC

The Transformation Gap: When Enterprise AI Investment Meets Organizational Reality

Our research reveals why generative AI isn't delivering



The funding is there.
What's missing: **Organizational Transformation**

Four focus areas
to close the gap

ESTABLISH TRUST

45%
of leaders are stuck
in pilot purgatory

Pilots don't scale without trust. Trust requires knowing your system's accuracy, failure points, and improvement paths. What doesn't get trusted doesn't get used—or create value.

Original Research: We are writing the playbook on evals:

Merkle is setting the bar on evaluating trust for agentic analytics applications.

01

UNIFIED EVAL FRAMEWORK

The first published framework evaluating both SQL & agent behavior

02

PROVEN METHODOLOGY

Methodology behind our eval app, validated on client implementations

03

OPEN PUBLICATION

Publishing openly on arXiv mid-March to establish and raise standards

COMING SOON

arXIV Paper

EvalForge: A Unified Evaluation Framework for Text-to-SQL and Agentic AI Systems in Enterprise Settings

Shannon Serrao^{*} Soumitra Chatterjee^{*} Nathan Miller^{*}
February 23, 2026

Abstract

Evaluating enterprise AI systems that translate natural language into SQL—and that increasingly wrap this capability inside multi-step agentic reasoning pipelines—poses challenges that existing academic benchmarks were not designed to address. Spider and BIRD established rigorous component-matching and execution accuracy protocols on controlled schemas, while G-Eval and RAGAS advanced LLM-based quality assessment of generated text; yet no single framework unifies these disciplines for the realities of production deployment. We present EvalForge, developed at Merkle, as the first unified evaluation framework that coherently integrates Text-to-SQL assessment with agentic behaviour evaluation into a single, enterprise-grade pipeline. EvalForge’s key differentiating contributions are threefold. First, it introduces *LLM-assisted SQL component extraction* that extends the Spider methodology to handle the CTE-heavy, dialect-specific SQL produced by modern LLMs—a gap that breaks every deterministic parser tested on production queries. Second, it introduces a *hybrid deterministic execution accuracy* algorithm that resolves the column-aliasing and numeric-tolerance brittleness of standard execution accuracy by using an LLM to infer structural alignments before applying fully deterministic cell-level scoring—bridging the gap between LLM-based and execution-based evaluation. Third, it introduces an *enterprise agentic evaluation suite* that operationalises tool-call fidelity, response faithfulness, G-Eval summary quality, and intent resolution as first-class metrics alongside SQL correctness—metrics entirely absent from the Text-to-SQL literature. Together these contributions make EvalForge the most comprehensive evaluation system currently available for enterprise natural-language data agents, and we describe the design decisions, algorithms, and deployment lessons that underpin it.

Contents

1	Introduction	1
2	Background and Related Work	3
2.1	The Text to SQL Evaluation Landscape	3
2.2	LLM-Based Evaluation	3
2.3	Agentic AI Evaluation	4
2.4	Summary: Gaps Addressed by EvalForge	4
3	Framework Design	5
3.1	Design Principles	5
3.2	Evaluation Architecture	5
4	Text-to-SQL Evaluation Methodology	7
4.1	Systemic Validity	7
4.2	LLM-Assisted Component Extraction	7

^{*}Merkle, a dentou company. Correspondence: shannon.serrao@merkle.com

